

TREAP

Huarui Liu, Jingzhou Qiu, Zicheng He

What is a Treap?


Keys in sorted order like a **BST**

Priorities follow the **heap property**

- Randomized BST Combines BST and heap properties
- A type of Cartesian Tree
- lookup, insertion, and removal in $O(\log N)$

Additional Operations:

- Split: $O(\log N)$
- Merge: $O(\log N)$



Why Treap?

Self-Balancing via random priorities


Simpler to implement than AVL or Red-Black trees

Can be modified to support segment tree operations and even more— all in $O(\log N)$

- Reverse on the interval.
- Addition / painting on the interval.

Applications

- Linux kernel page cache management
- General Purpose Allocator (GPA)



Priority for Max/Min Heap
Key for Binary Search Tree

Insert (Min-heap)


- Pick a random priority/specify a priority
- Insert as inserting in BST
- Rotate until the heap order is maintained

Runtime: $O(\log N)$

```
function insert(node, key, priority):  
    if node is empty:  
        create and return a new node with key and priority  
  
    if key is less than node's key:  
        recursively insert into left subtree  
        if left child has higher priority than current node:  
            perform right rotation  
  
    else if key is greater than node's key:  
        recursively insert into right subtree  
        if right child has higher priority than current node:  
            perform left rotation  
  
    else:  
        // key is equal — duplicate, so do nothing  
  
    return current node
```

Insert-Example

Insert(15) -> Random priority=9



Delete(Min-heap)


- Find the node by key (BST-style).
- If the node has 0 or 1 child:
 - return non-null child or null
- If the node has 2 children:
 - Rotate the child with the smaller priority up
 - Recurse on the same key to delete it

Runtime: $O(\log N)$

```
function delete(node, key):  
    if node is null:  
        return null  
    if key < node.key:  
        node.left = delete(node.left, key)  
    else if key > node.key:  
        node.right = delete(node.right, key)  
    else:  
        if node has at most one child:  
            return the non-null child (or null)  
        if left.priority < right.priority:  
            rotate right, then delete key from right child  
        else:  
            rotate left, then delete key from left child  
    return node
```

Delete-Example

Delete (15)



Build

Builds a tree from a list of values.

Heapify ensures the parent node has the highest/lowest priority by recursively swapping with the larger/smaller-priority child

Case 1: Input Keys Are Sorted -> Build in $O(N)$ time

Select the middle element to construct BST


Use heapify to ensure the heap property based on priorities



Case 2: Perform N insertions -> $O(N \log N)$ time

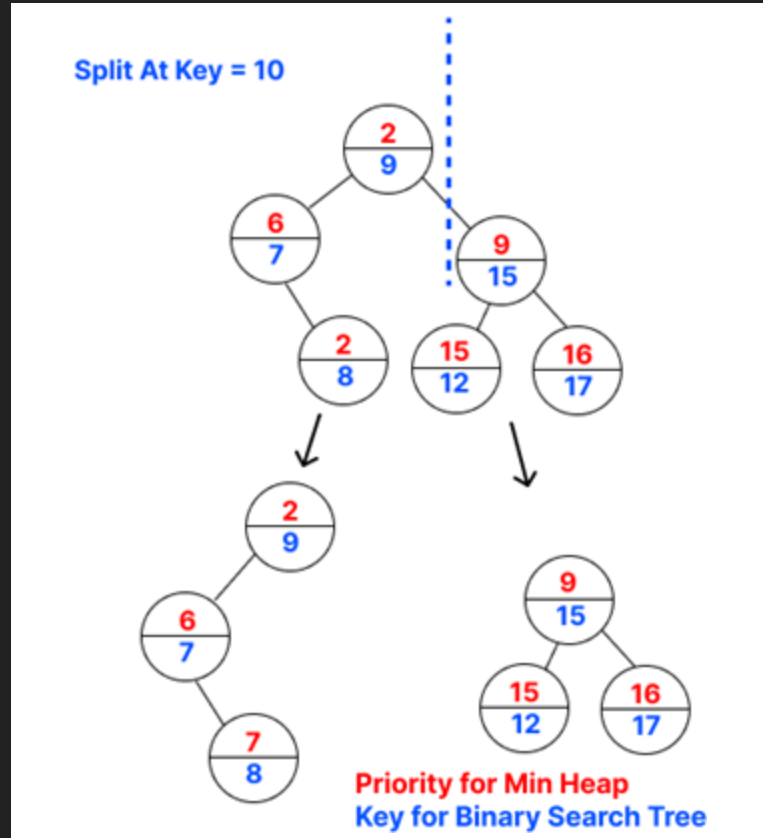
Split

- Decide which subtree the root node would belong to (left or right)
- Recursively call split on one of its children
- Create the final result by reusing the recursive split call
- Runtime: $O(\log N)$



```
struct SplitNodes { Node* left; Node* right; };  
function: split(node, key)  
    If node is null:  
        return (null, null)  
    If key <= node.key:  
        (left, right) = split(node.left, key)  
        node.left = right  
        return (left, node)  
    Else:  
        (left, right) = split(node.right, key)  
        node.right = left  
        return (node, right)
```


Split-Example



Merge(Min-heap)

- Merges two treaps (left and right) assuming all keys in left are less than those in right.
- Chooses the root with larger/smaller priority to maintain the heap property

Runtime: $O(\log N)$



```
function: merge(left, right)
```

```
  If left is null or right is null:
```

```
    return left if left exists, otherwise right
```

```
  If left.priority < right.priority:
```

```
    left.right = merge(left.right, right)
```


```
    return left
```

```
  Else:
```

```
    right.left = merge(left, right.left)
```

```
    return right
```

Merge-Example




Insert-Using Split

- Pick a random priority/specify a priority
- Insert as inserting in BST
- Rotate until the heap order is maintained

Runtime: $O(\log N)$

```
function insertSplit(key, optional_priority):  
    if optional_priority is not provided:  
        priority = generate_random_priority()  
    else:  
        priority = optional_priority  
    new_node = new Node(key, priority)  
  
    (L, R) = split(root, key)  
    root = merge(merge(L, new_node), R)  
  
    update_subtree_size(root)
```



Delete-Using Split


- The tree is split into three parts:
 - L: nodes with keys $< \text{key}$
 - target: node with key $= \text{key}$
 - R: nodes with keys $> \text{key}$
- Delete the target node
- Merge L and R back together

Runtime: $O(\log N)$

```
function deleteSplit(key):  
    (L, mid) = split(root, key)  
    (target, R) = split(mid, key + 1)
```

```
    delete target  
    root = merge(L, R)
```

```
    update_subtree_size(root)
```



Extension on Treaps-Implicit Treap

Implicit Keys:

- The position of each node is its key
- Index is calculated using subtree sizes

Nodes store more info to support extra features

Allow fast range queries and updates

Flexible Operations: Allows insertions, deletions, and reversals

```
struct TreapNode {  
    int value;  
    int priority;  
    int subtreeSize;  
    bool reversed;  
    int addValue;  
    bool needsUpdate;  
  
    TreapNode* leftChild;  
    TreapNode* rightChild;  
}
```

Utility functions

```
function count(node):  
    if node is null:  
        return 0  
    return node.subtreeSize  
// subtreeSize: subtree node count including current node
```

```
function updateCount(node):  
    if node is not null:  
        node.subtreeSize = 1 + count(node.left) + count(node.right)
```

SubtreeSize serves as the node's position like an **implicit** index

updateCount function maintains accurate subtree sizes

We will implement an Implicit Treap with the reverse function

Updated Split

Runtime: $O(\log N)$

```
function split(node, index):  
    if node is null:  
        return (null, null)  
  
    push(currentNode)  
    compute currentNodeIndex based on size of left subtree  
  
    if index <= currentNodeIndex:  
        split left subtree  
        attach result to node's left  
        return (left part, node)  
    else:  
        split right subtree  
        attach result to node's right  
        return (node, right part)  
  
    update node's metadata (e.g., subtree size)
```

- calculates each node's position
- split the nodes with into left and right subtree based on the given index
- uses push() to apply any pending lazy updates (in our example, the reversal flag only)
- use updateCount() to keep track of subtree sizes accurate

*push() function will be introduced later

Updated Merge

```
function merge(left, right):  
    push(leftTree)  
    push(rightTree)  
  
    if either tree is null:  
        return the non-null tree  
  
    if left.priority < right.priority:  
        left.right = merge(left.right, right)  
        update left's metadata  
        return left  
    else:  
        right.left = merge(left, right.left)  
        update right's metadata  
        return right
```

- uses push() to apply any pending lazy updates
- same merging logic, except we need to keep track of subtree sizes and other metadata

Runtime: $O(\log N)$

Lazy Propagation for Reverse

```
function push(node):  
    if node is not null and node.reversed is true:  
        swap(node.left, node.right)  
  
    if node.left is not null:  
        node.left.reversed = not node.left.reversed  
  
    if node.right is not null:  
        node.right.reversed = not node.right.reversed  
  
    node.reversed = false
```

- Instead of immediately swapping all left/right children(which would be $O(N)$), flip the reversed flag
- Call push() when the Node is accessed

Runtime: $O(\log N)$

Reverse Function

Runtime: $O(\log N)$

```
function reverseSegment(root, leftIndex, rightIndex):  
    // Split the tree into two parts:  
    (leftSubtree, remainingTree) = split(root, leftIndex)  
  
    // Split remainingTree again:  
    (middleSubtree, rightSubtree) = split(remainingTree, rightIndex - leftIndex + 1)  
  
    // Mark the middleSubtree to be reversed (lazy propagation)  
    if middleSubtree is not null:  
        middleSubtree.reverseFlag = not middleSubtree.reverseFlag  
  
    // Merge all three parts back together in order  
    root = merge(leftSubtree, middleSubtree)  
    root = merge(root, rightSubtree)
```

Conclusion

- Treap combines the balanced BSTs with array-style indexing with most operations in $O(\log N)$ time
- more powerful than segment trees

However:

- They're mostly only used in competitive programming and academia:
- developers favor other more commonly used structures like segment trees or balanced BST libraries

Reference

https://cp-algorithms.com/data_structures/treap.html

<https://www.youtube.com/watch?v=6x0UIIBLRsc>

<https://courses.cs.washington.edu/courses/cse326/00wi/handouts/lecture19/sld017.htm>